

Food Desert Code

Kaitlyn Keebaugh, Matthew Pergolski, Miranda Braman, Victor Yamaykin
IST707

Section 1: Rule Mining

LIBRARIES

```
library(knitr)
library(magrittr)
library(tidyverse)
library(DataExplorer)
library(ggplot2)
library(arules)
library(arulesViz)
library(ROSE)
library(randomForest)
library(caret)
library(e1071)
library(dplyr)
library(RColorBrewer)
library(neuralnet)
library(arm)
library(rpart.plot)
library(rpart)
library(rattle)
library(plotly)
library(tidyverse)
library(cluster)
library(factoextra)
library(dendextend)
library(DescTools)
```

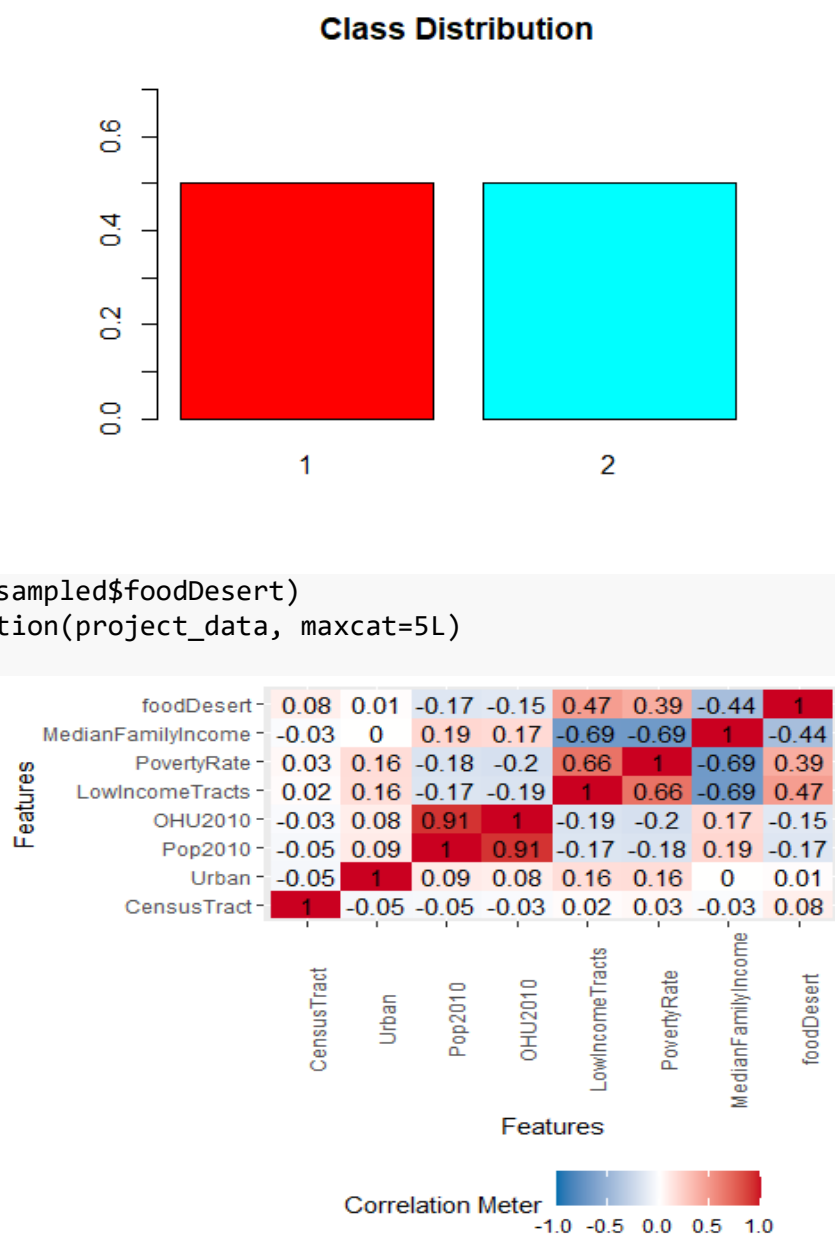
LOAD PACKAGES & DATA SET

```
# Load Data Set
project_data <- read.csv("C:\\Users\\kkeep\\Documents\\Balanced_FoodDeser
tComparisonData_2006_2019.csv",stringsAsFactors = TRUE)

introduce(project_data)
table(project_data$foodDesert)

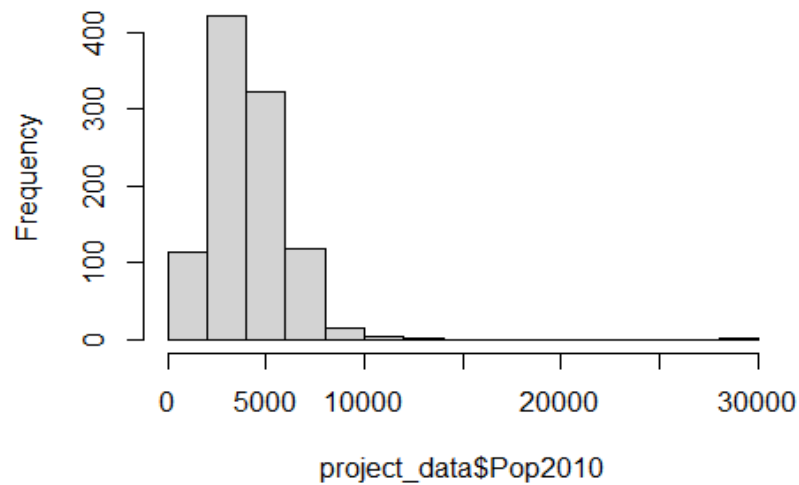
# Oversampling test
set.seed(42)
barplot(prop.table(table(project_data$foodDesert)),
        col = rainbow(2),
        ylim = c(0,0.7),
        main = "Class Distribution")
```

```
under_sampled <- ovun.sample(foodDesert~., data=project_data, seed=1, method="under")$data
```



```
table(discretize(project_data$Pop2010, breaks=3))
hist(project_data$Pop2010, breaks = 12, main = "Equal Frequency Discretization")
```

Equal Frequency Discretization




```
##
##      [36,3.13e+03) [3.13e+03,4.73e+03) [4.73e+03,2.97e+04]
##      333          333          334

ARM_ready_data <- discretizeDF(project_data, methods = list(
  Pop2010 = list(method = "frequency", breaks = 5,
    labels = c("very low", "low", "medium", "high", "very high")
  ),
  OHU2010 = list(method = "frequency", breaks = 3,
    labels = c("low", "medium", "high")),
  PovertyRate = list(method = "frequency", breaks = 3,
    labels = c("low", "medium", "high")),
  MedianFamilyIncome = list(method = "frequency", breaks = 3,
    labels = c("low", "medium", "high"))), default = 1
ist(method = "none"))

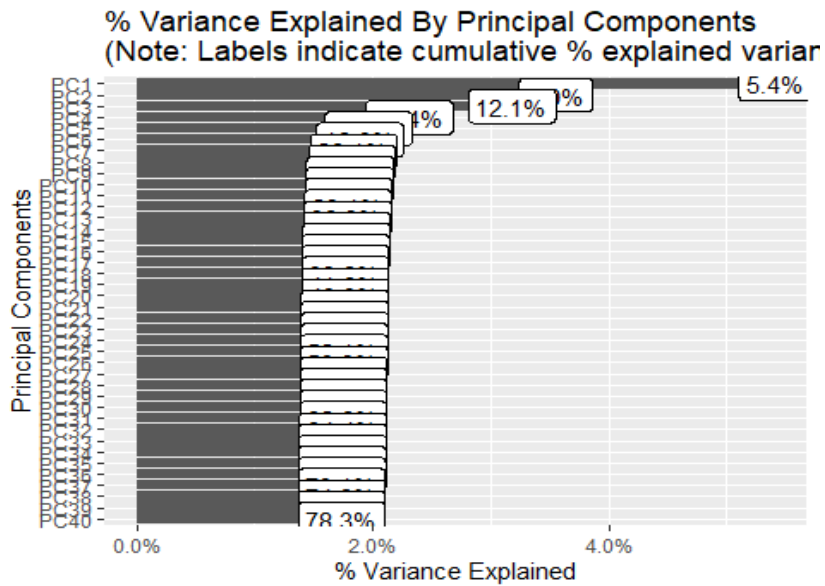
# Turn the data set to transactions
tid <- as.character(ARM_ready_data$CensusTract)
ARM_ready_data$id <- NULL
transactions <- as(ARM_ready_data, "transactions")
transactionInfo(transactions)[["transactionID"]] <- tid

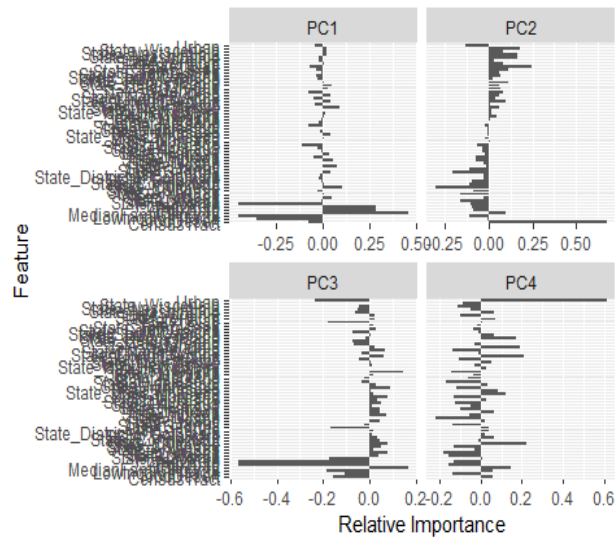
str(project_data)
plot_str(project_data)
introduce(project_data)

plot_correlation(project_data, maxcat=5L)
```




```
plot_prcomp(na.omit(project_data), variance_cap = 0.8, nrow = 2L, ncol = 2L)
```







Page 1




Page 2




Page 3




Page 4




Page 5




Page 6




Page 7



Page 8



Page 9

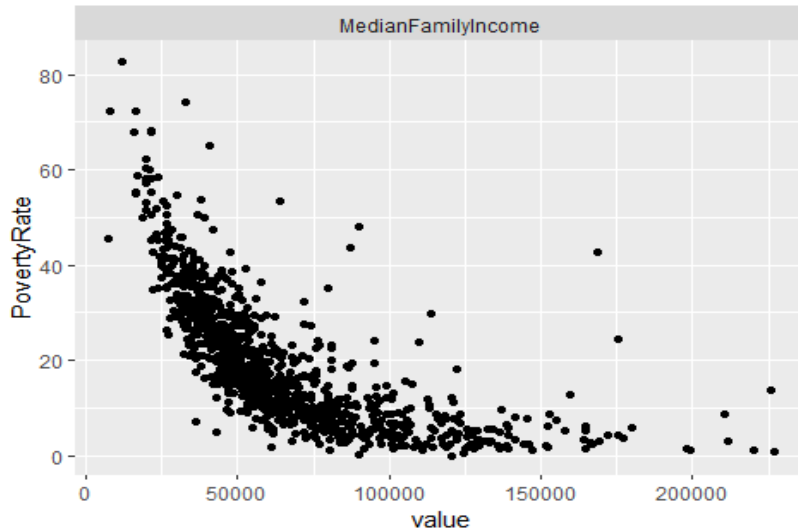


Page 10

```
plot(project_data$PovertyRate, project_data$MedianFamilyIncome)
```



```
plot_scatterplot(project_data[, c("PovertyRate", "MedianFamilyIncome")], by="PovertyRate", sampled_rows = 1000L)
```



```
p <- ggplot(project_data[sample(nrow(project_data), 250), ], aes(MedianFamilyIncome, PovertyRate, color=foodDesert, size = Pop2010*10,
                                                                    main="Food"))
) + geom_point(na.rm = T)
```

##	CensusTract	State	County	Urban	Pop2010	OHU2010
## 1	17031381800	Illinois	Cook County	1	very low	low
## 2	42003483800	Pennsylvania	Allegheny County	1	low	medium
## 3	13117130612	Georgia	Forsyth County	1	very high	medium
## 4	45091060905	South Carolina	York County	1	very high	high
## 5	6037501504	California	Los Angeles County	1	medium	medium
## 6	48141010336	Texas	El Paso County	1	very high	high

##	LowIncomeTracts	PovertyRate	MedianFamilyIncome	foodDesert
## 1	1	medium	medium	1
## 2	1	high	low	1
## 3	0	low	high	1
## 4	0	medium	medium	1
## 5	1	high	low	1
## 6	0	medium	medium	1

```
## 'data.frame': 1000 obs. of 10 variables:
## $ CensusTract : num 1.70e+10 4.20e+10 1.31e+10 4.51e+10 6.04e+09 .
## $ State : chr "Illinois" "Pennsylvania" "Georgia" "South Carolina" ...
## $ County : chr "Cook County" "Allegheny County" "Forsyth County" "York County" ...
## $ Urban : int 1 1 1 1 1 1 0 1 1 1 ...
## $ Pop2010 : int 1188 3165 5614 8916 3592 6638 2128 1193 2632 2300 ...
## $ OHU2010 : int 591 1485 1740 3663 1588 1986 822 606 815 1066 ...
## $ LowIncomeTracts : int 1 1 0 0 1 0 0 1 1 1 ...
## $ PovertyRate : num 23.4 35.2 1.8 17.3 24.3 13.7 10.7 41.9 33 20.4
```

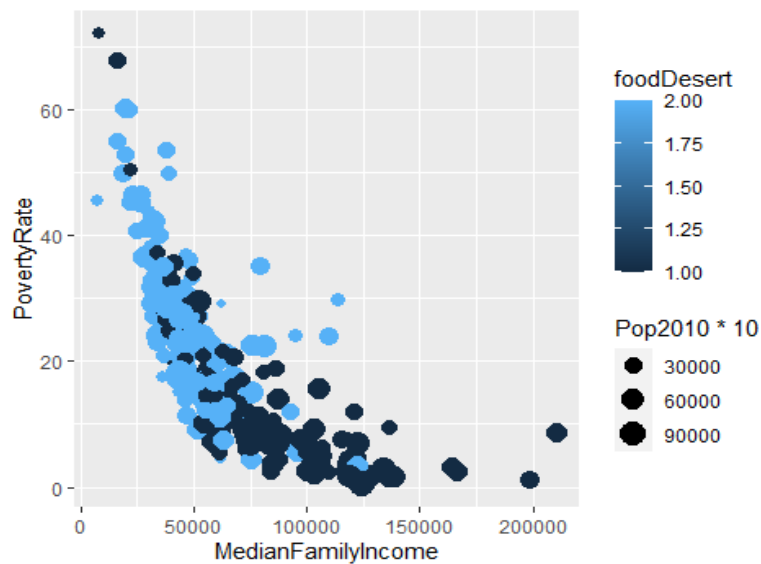


```

...
## $ MedianFamilyIncome: int  60938 26336 151944 60625 45208 56510 85703 254
41 39958 56797 ...
## $ foodDesert          : int  1 1 1 1 1 1 1 1 1 1 ...

## rows columns discrete_columns continuous_columns all_missing_columns
## 1 1000      10              2              8              0
## total_missing_values complete_rows total_observations memory_usage
## 1              0              1000              10000      93816

```




Look at relative frequency plot to see how many times these items have appeared as compared to others

```

itemFrequencyPlot(transactions, topN=20, type="relative", col=brewer.pal(8,
'Pastel2'),main="Relative Item Frequency Plot for Food Access Research Atlas
(FARA)")

```

Relative Item Frequency Plot for Food Access Research At



APRIORI

```
## Apriori
# Get the rules with low support and low confidence
rules <- apriori(transactions, parameter = list(supp = 0.02, conf = 0.7, minl
en=4))

# Show rules
inspect(rules[1:20])

## Sort by lift
SortedRules_conf <- sort(rules, by="confidence", decreasing=F)
inspect(SortedRules_conf[1:50])

## Take the top 10 rules sorted by lift
top10rules_conf <- head(SortedRules_conf, n = 10, by = "confidence")
inspect(top10rules_conf)

## Parameter specification:
## confidence minval smax arem aval originalSupport maxtime support minlen
## 0.7 0.1 1 none FALSE TRUE 5 0.02 4
## maxlen target ext
## 10 rules TRUE
##
## Algorithmic control:
## filter tree heap memopt load sort verbose
## 0.1 TRUE TRUE FALSE TRUE 2 TRUE
##
## Absolute minimum support count: 20
##
## set item appearances ...[0 item(s)] done [0.00s].
## set transactions ...[70 item(s), 1000 transaction(s)] done [0.00s].
## sorting and recoding items ... [41 item(s)] done [0.00s].
## creating transaction tree ... done [0.00s].
## checking subsets of size 1 2 3 4 5 6 7 8 done [0.00s].
## writing ... [7915 rule(s)] done [0.00s].
## creating S4 object ... done [0.00s].

## lhs rhs
support confidence coverage lift count
## [1] {CensusTract=[3.72e+10,5.51e+10],
## State=Tennessee,
## Urban=[0,1]} => {LowIncomeTracts=[0,1]}
0.02 1 0.02 1 20
## [2] {CensusTract=[3.72e+10,5.51e+10],
## State=Tennessee,
## LowIncomeTracts=[0,1]} => {Urban=[0,1]}
0.02 1 0.02 1 20
## [3] {State=Tennessee,
## Urban=[0,1],
```

```

##      LowIncomeTracts=[0,1]}      => {CensusTract=[3.72e+10,5.51e+10
]}  0.02      1      0.02      3      20
## [4] {CensusTract=[3.72e+10,5.51e+10]
,
##      State=Tennessee,
##      Urban=[0,1]}      => {foodDesert=[1,2]}
0.02      1      0.02      1      20
## [5] {CensusTract=[3.72e+10,5.51e+10],
##      State=Tennessee,
##      foodDesert=[1,2]}      => {Urban=[0,1]}
0.02      1      0.02      1      20
## [6] {State=Tennessee,
##      Urban=[0,1],
##      foodDesert=[1,2]}      => {CensusTract=[3.72e+10,5.51e+10
]}  0.02      1      0.02      3      20
## [7] {CensusTract=[3.72e+10,5.51e+10],
##      State=Tennessee,
##      LowIncomeTracts=[0,1]}      => {foodDesert=[1,2]}
0.02      1      0.02      1      20
## [8] {CensusTract=[3.72e+10,5.51e+10],
##      State=Tennessee,
##      foodDesert=[1,2]}      => {LowIncomeTracts=[0,1]}
0.02      1      0.02      1      20
## [9] {State=Tennessee
,
##      LowIncomeTracts=[0,1],
##      foodDesert=[1,2]}      => {CensusTract=[3.72e+10,5.51e+10
]}  0.02      1      0.02      3      20
## [10] {State=Tennessee,
##      Urban=[0,1],
##      LowIncomeTracts=[0,1]}      => {foodDesert=[1,2]}

##      lhs      rhs      sup
port confidence coverage lift count
## [1] {CensusTract=[1.81e+10,3.72e+10),
##      Urban=[0,1],
##      PovertyRate=medium}      => {MedianFamilyIncome=medium}      0
.070      0.70      0.100      2.1      70
## [2] {CensusTract=[1.81e+10,3.72e+10),
##      LowIncomeTracts=[0,1],
##      PovertyRate=medium}      => {MedianFamilyIncome=medium}      0
.070      0.70      0.100      2.1      70
## [3] {CensusTract=[1.81e+10,3.72e+10),
##      PovertyRate=medium,
##      foodDesert=[1,2]}      => {MedianFamilyIncome=medium}      0
.070      0.70      0.100      2.1      70
## [4] {CensusTract=[1.81e+10,3.72e+10),
##      Urban=[0,1],
##      LowIncomeTracts=[0,1],

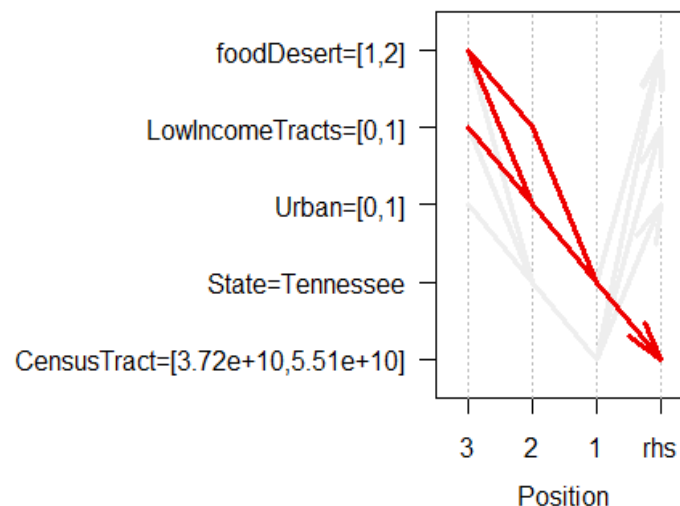
```

```
##      PovertyRate=medium}          => {MedianFamilyIncome=medium}    0
.070      0.70      0.100  2.1      70

##      lhs                                rhs
support confidence coverage lift count
## [1] {CensusTract=[3.72e+10,5.51e+10],
##      State=Tennessee,
##      Urban=[0,1]}          => {LowIncomeTracts=[0,1]}
0.02      1      0.02      1      20
## [2] {CensusTract=[3.72e+10,5.51e+10],
##      State=Tennessee,
##      LowIncomeTracts=[0,1]}    => {Urban=[0,1]}
0.02      1      0.02      1      20
## [3] {State=Tennessee,
##      Urban=[0,1],
##      LowIncomeTracts=[0,1]}    =>

## Visualize the rules with a parallel coordinate plot
plot(top10rules_conf, method = "paracoord")
plot(top10rules_conf, method = "graph", interactive = T)
```

Parallel coordinates plot for 10 rules



Section 2: Clustering

PRE-PROCESSING / DATA MUNGING

```
project_data -> raw.data.cluster

str(raw.data.cluster)

## 'data.frame':   1000 obs. of  10 variables:
## $ CensusTract      : num  1.70e+10 4.20e+10 1.31e+10 4.51e+10 6.04e+09 .
## ..
```

```

## $ State          : chr  "Illinois" "Pennsylvania" "Georgia" "South Carolina" ...
## $ County         : chr  "Cook County" "Allegheny County" "Forsyth County" "York County" ...
## $ Urban          : int   1 1 1 1 1 1 0 1 1 1 ...
## $ Pop2010        : int   1188 3165 5614 8916 3592 6638 2128 1193 2632 2300 ...
## $ OHU2010        : int    591 1485 1740 3663 1588 1986 822 606 815 1066 ...
## $ LowIncomeTracts : int    1 1 0 0 1 0 0 1 1 1 ...
## $ PovertyRate     : num   23.4 35.2 1.8 17.3 24.3 13.7 10.7 41.9 33 20.4 ...
## $ MedianFamilyIncome: int   60938 26336 151944 60625 45208 56510 85703 25441 39958 56797 ...
## $ foodDesert      : int    1 1 1 1 1 1 1 1 1 1 ...

## rows columns discrete_columns continuous_columns all_missing_columns
## 1 1000      10      2      8      0
## total_missing_values complete_rows total_observations memory_usage
## 1      0      1000      10000      93816

##
## 1 2
## 500 500

# Convert Data to Numeric-Only For Clustering
colnames(data.cluster[,c(1,2,3,4)])

## [1] "Group.1" "CensusTract" "State" "County"

str(data.cluster[,c(-1,-2,-3,-4)])

## 'data.frame': 50 obs. of 7 variables:
## $ Urban : num 0.667 0 0.833 0.556 0.901 ...
## $ Pop2010 : num 4060 4193 4378 3422 4608 ...
## $ OHU2010 : num 1604 1688 1573 1375 1563 ...
## $ LowIncomeTracts : num 0.524 0 0.583 0.778 0.582 ...
## $ PovertyRate : num 24.7 9.8 20.9 22.1 15.9 ...
## $ MedianFamilyIncome: num 50854 92000 56345 50307 79064 ...
## $ foodDesert : num 1.48 1 1.46 1.78 1.32 ...

num.data.cluster <- data.cluster[,c(-1,-2,-3,-4)]
num.data.cluster <- as.data.frame(scale(num.data.cluster))
str(num.data.cluster)


## 'data.frame': 50 obs. of 7 variables:
## $ Urban : num -0.128 -2.877 0.559 -0.586 0.838 ...
## $ Pop2010 : num 0.214 0.408 0.678 -0.719 1.014 ...
## $ OHU2010 : num 0.422 0.768 0.293 -0.531 0.249 ...
## $ LowIncomeTracts : num -0.3397 -2.5871 -0.0843 0.75 -0.0882 ...
## $ PovertyRate : num 0.898 -1.585 0.268 0.478 -0.557 ...

```

```
## $ MedianFamilyIncome: num -0.728 1.793 -0.391 -0.761 1 ...
## $ foodDesert          : num -0.133 -2.289 -0.214 1.232 -0.846 ...
```

MODELS

```
# Optimal Amount of Clusters | Average Silhouette Method
fviz_nbclust(num.data.cluster, FUN = hcut, method = "silhouette")
# plot shows 3 optimal clusters
```



Agnes Function | Dendrogram

```
# Determine Optimal Agnes Method
m.assess <- c("average", "single", "complete", "ward")
names(m.assess) <- c("average", "single", "complete", "ward")

compute.coeff <- function(x)
{
  agnes(num.data.cluster, method = x)$ac
}

hc.coeff.df <- as.data.frame(map_dbl(m.assess, compute.coeff))
hc.coeff.df

##           map_dbl(m.assess, compute.coeff)
## average                                0.7337589
## single                                 0.5997242
```

```
## complete                0.8469888
## ward                    0.8895291

# Method 'ward' conveys highest quality with 0.8895291


H.C <- agnes(num.data.cluster, method = "ward")

# Agglomerative coefficient (which measures the amount of clustering structure found)
# (values closer to 1 suggest strong clustering structure)
H.C$ac

## [1] 0.8895291

## [1] 0.8895291

pltree(H.C, cex = 0.6, hang = -1, main = "Dendrogram of AGNES")
rect.hclust(H.C, k = 3, border = 2:5)
```




```
num.data.cluster
agnes (*, "ward")
```

hclust Function / Dendrogram

```
# Dissimilarity matrix
d <- dist(num.data.cluster, method = "euclidean")
#d
# Hierarchical clustering using Complete Linkage
hc1 <- hclust(d, method = "ward.D" )

# Plot the obtained dendrogram
```

```
plot(hc1, cex = 0.6, hang = -1)
rect.hclust(hc1, k = 3, border = 2:5)
```



```
#fviz_nbclust(num.data.cluster, kmeans, method = "wss")
```

K MEANS

```
k.means <- kmeans(num.data.cluster, 3)
k.means
```

```
## K-means clustering with 3 clusters of sizes 4, 15, 31
##
## Cluster means:
##      Urban      Pop2010      OHU2010 LowIncomeTracts PovertyRate
## 1 -0.9086002  0.8152088  0.7177478      -2.26061986  -1.7716848
## 2 -0.2280598 -0.9927497 -1.0534376       0.78546056   0.9185914
## 3  0.2275903  0.3751745  0.4171152      -0.08836867  -0.2158752
##      MedianFamilyIncome foodDesert
## 1          2.2236455 -1.2306823
## 2         -0.8477377  0.4461845
## 3          0.1232736 -0.0570980
##
## Clustering vector:
##      Alabama      Alaska      Arizona
##           3           1           3
##      Arkansas      California      Colorado
##           2           3           3
##      Connecticut      Delaware District of Columbia
##           3           2           2
##           Florida      Georgia      Hawaii
##           3           3           1
```



```

##          Idaho          Illinois          Indiana
##          3            3            2
##          Iowa          Kansas          Kentucky
##          3            3            2
##          Louisiana      Maine          Maryland
##          2            2            3
##          Massachusetts  Michigan      Minnesota
##          3            3            3
##          Mississippi    Missouri      Montana
##          2            3            3
##          Nebraska        Nevada        New Hampshire
##          2            3            3
##          New Jersey      New Mexico      New York
##          1            2            3
##          North Carolina  North Dakota    Ohio
##          2            1            2
##          Oklahoma        Oregon          Pennsylvania
##          3            3            3
##          Rhode Island    South Carolina  South Dakota
##          2            3            3
##          Tennessee        Texas          Utah
##          2            3            3
##          Vermont          Virginia      Washington
##          2            3            3
##          West Virginia    Wisconsin
##          3            3
##
## Within cluster sum of squares by cluster:
## [1] 20.25762 90.21556 84.16343
## (between_SS / total_SS = 43.3 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withi
nss"
## [6] "betweenss"    "size"         "iter"         "ifault"

k.means$centers

##          Urban    Pop2010    OHU2010 LowIncomeTracts PovertyRate
## 1 -0.9086002  0.8152088  0.7177478    -2.26061986  -1.7716848
## 2 -0.2280598 -0.9927497 -1.0534376     0.78546056   0.9185914
## 3  0.2275903  0.3751745  0.4171152    -0.08836867  -0.2158752
## MedianFamilyIncome foodDesert
## 1          2.2236455 -1.2306823
## 2          -0.8477377  0.4461845
## 3           0.1232736 -0.0570980

assignment_clusters <- data.frame(num.data.cluster, k.means$cluster)
assignment_clusters


```

```
head(assignment_clusters)
```

```
##           Urban    Pop2010    OHU2010 LowIncomeTracts PovertyRate
## Alabama   -0.1280735  0.2137247  0.4222219   -0.33966663   0.8983645
## Alaska    -2.8767226  0.4084355  0.7684432   -2.58707088  -1.5848068
## Arizona    0.5590888  0.6782671  0.2926119   -0.08427978   0.2676231
## Arkansas  -0.5861817 -0.7186934 -0.5309601    0.74998392   0.4784008
## California 0.8384845  1.0141702  0.2494695   -0.08820881  -0.5570090
## Colorado   0.5952552 -0.0234643  0.1716279   -0.10309776  -0.5221703
##           MedianFamilyIncome foodDesert k.means.cluster
## Alabama      -0.7276807  -0.1329249                3
## Alaska        1.7928665  -2.2886065                1
## Arizona       -0.3912764  -0.2137629                3
## Arkansas      -0.7611959   1.2323401                2
## California     1.0004205  -0.8459580                3
## Colorado       0.3378109   1.0470272                3
```


```
clusplot(num.data.cluster, k.means$cluster, color=T, shade=T,
         Labels=2, lines=0) # plot clusters
```

CLUSPLOT(num.data.cluster)



These two components explain 73.25 % of the point variab

```
fviz_cluster(k.means, data = num.data.cluster,
             main = 'Cluster Plot',
             xlab = '',
             ylab = '', pointsize = num.data.cluster$PovertyRate)
```



```
# Helpful Data Tables
```

```
# Create Separate DF
```

```
main.cluster.df <- data.frame(raw.data.cluster, k.means$cluster)
```

```
## Warning in data.frame(raw.data.cluster, k.means$cluster): row names were found
```

```
## from a short variable and have been discarded
```

```
main.cluster.df <- main.cluster.df[,c(-1,-3,-4)]
head(main.cluster.df)
```

```
##           State Pop2010 OHU2010 LowIncomeTracts PovertyRate MedianFamilyIncome
## 1      Illinois    1188     591             1         23.4
60938
## 2  Pennsylvania    3165    1485             1         35.2
26336
## 3      Georgia     5614    1740             0          1.8
51944
## 4 South Carolina     8916    3663             0         17.3
60625
## 5    California     3592    1588             1         24.3
45208
## 6      Texas       6638    1986             0         13.7
56510
## foodDesert k.means.cluster
## 1          1              3
## 2          1              1
## 3          1              3
## 4          1              2
## 5          1              3
## 6          1              3
```

```

# Discretize Poverty Rate
pv.bins <- 3

min.pv <- min(main.cluster.df$PovertyRate)
min.pv

## [1] 0

max.pv <- max(main.cluster.df$PovertyRate)
max.pv

## [1] 82.8

mid.pv <- (max.pv - min.pv) / pv.bins
mid.pv

## [1] 27.6

mid.pv * 3

## [1] 82.8

main.cluster.df$DiscPovertyRate <- cut(main.cluster.df$PovertyRate,
                                       breaks = c(min.pv, mid.pv, max.pv,
Inf),
                                       labels = c('Min', 'Mid', 'Max'))

str(main.cluster.df)

## 'data.frame':    1000 obs. of  9 variables:
## $ State          : Factor w/ 50 levels "Alabama","Alaska",...: 14 39 11
41 5 44 28 39 11 10 ...
## $ Pop2010        : int  1188 3165 5614 8916 3592 6638 2128 1193 2632 2
300 ...
## $ OHU2010        : int  591 1485 1740 3663 1588 1986 822 606 815 1066
...
## $ LowIncomeTracts : int  1 1 0 0 1 0 0 1 1 1 ...
## $ PovertyRate     : num  23.4 35.2 1.8 17.3 24.3 13.7 10.7 41.9 33 20.4
...
## $ MedianFamilyIncome: int  60938 26336 151944 60625 45208 56510 85703 254
41 39958 56797 ...
## $ foodDesert      : int  1 1 1 1 1 1 1 1 1 1 ...
## $ k.means.cluster  : int  3 1 3 2 3 3 3 2 2 3 ...
## $ DiscPovertyRate  : Factor w/ 3 levels "Min","Mid","Max": 1 2 1 1 1 1 1
2 2 1 ...

head(main.cluster.df)

##           State Pop2010 OHU2010 LowIncomeTracts PovertyRate MedianFamilyI
ncome
## 1      Illinois   1188     591                1        23.4
60938

```

```

## 2    Pennsylvania    3165    1485            1        35.2
26336
## 3          Georgia    5614    1740            0         1.8            1
51944
## 4 South Carolina    8916    3663            0        17.3
60625
## 5      California    3592    1588            1        24.3
45208
## 6          Texas    6638    1986            0        13.7
56510
##   foodDesert k.means.cluster DiscPovertyRate
## 1          1              3          Min
## 2          1              1          Mid
## 3          1              3          Min
## 4          1              2          Min
## 5          1              3          Min
## 6          1              3          Min

  cluster.1.df <- main.cluster.df[main.cluster.df$k.means.cluster == 1,]
  Mode(cluster.1.df$DiscPovertyRate)

## [1] Min
## attr(,"freq")
## [1] 59
## Levels: Min Mid Max

  cluster.1.df

  cluster.2.df <- main.cluster.df[main.cluster.df$k.means.cluster == 2,]
  Mode(cluster.2.df$DiscPovertyRate)

## [1] Min
## attr(,"freq")
## [1] 216
## Levels: Min Mid Max

  mean(cluster.2.df$PovertyRate)

## [1] 20.80967

  cluster.2.df

  cluster.3.df <- main.cluster.df[main.cluster.df$k.means.cluster == 3,]
  cluster.3.df <- na.omit(cluster.3.df)
  Mode(cluster.3.df$DiscPovertyRate)

## [1] Min
## attr(,"freq")
## [1] 472
## Levels: Min Mid Max

  cluster.3.df

```

```
main.cluster.df <- aggregate(main.cluster.df, by = list(main.cluster.df$State), FUN = mean)
main.cluster.df
```

Section 3: Classification Models

DECISION TREE PREP

*## Must normalize some columns and make their numbers between 0 and 1 for
##Classification Models*

```
set.seed(341)
```

#randomize the dataset

```
fooddesert[sample(nrow(fooddesert)),]-> fooddesert
```

#make train 80% of data and test 20%

```
nrow(fooddesert)*.8-> index
```

```
fooddesert[1:index,]->train
```

```
fooddesert[(index+1): nrow(fooddesert),]->test
```

check percentages

```
prop.table(table(train$foodDesert))
```

```
##
```

```
##      0      1
```

```
## 0.5075 0.4925
```

```
prop.table(table(test$foodDesert))
```

```
##
```

```
##      0      1
```

```
## 0.47 0.53
```

DECISION TREE MODELS


#decision tree 1

```
train_tree1 <- rpart(foodDesert ~ ., data = train,
                     method="class", control=rpart.control(cp=0, maxdepth=7))
```

```
predicted1= predict(train_tree1, test, type="class")
```

```
fancyRpartPlot(train_tree1)
```

```
table(FoodDesert=predicted1, true=data_test$foodDesert)
```




Decision tree #2

```
train_tree1 <- rpart(foodDesert ~
  Urban+OHU2010+LowIncomeTracts+PovertyRate+MedianFamilyIncome,
  data = train, method="class",
  control=rpart.control(cp=0.013, maxdepth=4))
```

#verify CP and size of tree

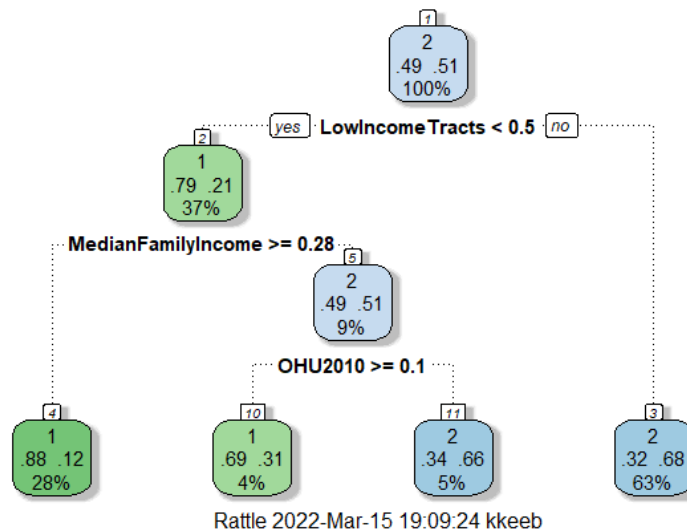
```
rsq.rpart(train_tree1)
plotcp(train_tree1)
```



```
printcp(train_tree1)
## Classification tree:
## rpart(formula = foodDesert ~ Urban + OHU2010 + LowIncomeTracts +
##   PovertyRate + MedianFamilyIncome, data = train, method = "class",
##   control = rpart.control(cp = 0.013, maxdepth = 4))
##
## Variables actually used in tree construction:
## [1] MedianFamilyIncome
```

```
##
## Root node error: 397/800 = 0.49625
##
## n= 800
##
##      CP nsplit rel error  xerror    xstd
## 1 0.40302      0  1.00000 1.07053 0.035553
## 2 0.01300      1  0.59698 0.61965 0.032877

predicted1= predict(train_tree1, test, type="class")
fancyRpartPlot(train_tree1)
```




```
confusionMatrix(predicted1, as.factor(test$foodDesert))
## Confusion Matrix and Statistics
##
##      Reference
## Prediction  1  2
##      1  63  11
##      2  40  86
##
##      Accuracy : 0.745
##      95% CI : (0.6787, 0.8039)
##      No Information Rate : 0.515
##      P-Value [Acc > NIR] : 2.232e-11
##
##      Kappa : 0.4939
##
##      McNemar's Test P-Value : 8.826e-05
##
##      Sensitivity : 0.6117
##      Specificity : 0.8866
##      Pos Pred Value : 0.8514
##      Neg Pred Value : 0.6825
##      Prevalence : 0.5150
```



```
##      Detection Rate : 0.3150
##      Detection Prevalence : 0.3700
##      Balanced Accuracy : 0.7491
##
##      'Positive' Class : 1
```

Decision Tree #3

```
train_tree1 <- rpart(foodDesert ~ LowIncomeTracts+MedianFamilyIncome,
                     data = train, method="class",
                     control=rpart.control(cp=0.005, maxdepth=3))
rsq.rpart(train_tree1)
plotcp(train_tree1)
```



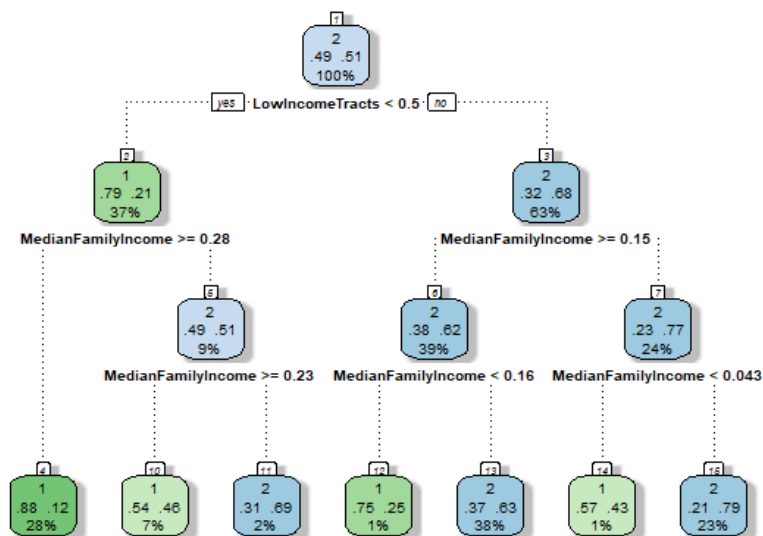
```
printcp(train_tree1)
## Classification tree:
## rpart(formula = foodDesert ~ LowIncomeTracts + MedianFamilyIncome,
##       data = train, method = "class", control = rpart.control(cp = 0.005,
##       maxdepth = 3))
##
## Variables actually used in tree construction:
## [1] LowIncomeTracts   MedianFamilyIncome
##
## Root node error: 397/800 = 0.49625
##
## n= 800
##
##      CP nsplit rel error xerror   xstd
## 1 0.4030227    0  1.00000 1.1058 0.035453
## 2 0.0062972    1  0.59698 0.6272 0.032987
## 3 0.0050000    3  0.58438 0.6272 0.032987

predicted1= predict(train_tree1, test, type="class")
fancyRpartPlot(train_tree1)
confusionMatrix(predicted1, as.factor(test$foodDesert))
```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  1  2
##           1 61  5
##           2 42 92
##
##           Accuracy : 0.765
##           95% CI : (0.7, 0.8219)
##           No Information Rate : 0.515
##           P-Value [Acc > NIR] : 3.060e-13
##
##           Kappa : 0.5347
##
## Mcnemar's Test P-Value : 1.512e-07
##
##           Sensitivity : 0.5922
##           Specificity : 0.9485
##           Pos Pred Value : 0.9242
##           Neg Pred Value : 0.6866
##           Prevalence : 0.5150
##           Detection Rate : 0.3050
##           Detection Prevalence : 0.3300
##           Balanced Accuracy : 0.7703
##
##           'Positive' Class : 1
printcp(train_tree1)

```



Rattle 2022-Mar-15 19:27:04 kkeeb

NEURAL NETWORKS

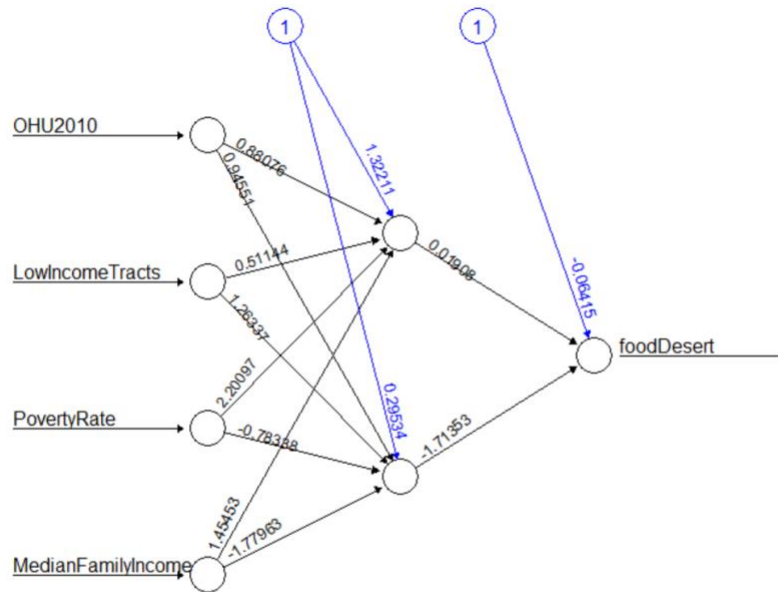
```
#set seed to get the same results each time  
set.seed(24)
```

```
#create a matrix with the columns you want to include in NN  
model.matrix(~OHU2010+LowIncomeTracts+PovertyRate+MedianFamilyIncome+foodDesert,  
             data=train)->train_matrix  
model.matrix(~OHU2010+LowIncomeTracts+PovertyRate+MedianFamilyIncome+foodDesert,  
             data=test)->test_matrix
```

```
#create formulas in order to run it through the neural network  
#the first is with training data.  
#we take out the first and last column otherwise we get an integer column  
#and another fooddesert column  
col_list <- paste(c(colnames(train_matrix[, -c(1,6)])), collapse="+")  
col_list <- paste(c("foodDesert~", col_list), collapse="")  
f <- formula(col_list)
```

```
#create a formula for the test matrix  
col_list <- paste(c(colnames(test_matrix[, -c(1,6)])), collapse="+")  
col_list <- paste(c("foodDesert~", col_list), collapse="")  
m <- formula(col_list)
```

```
#design and run the neural network  
neuralnet(f, data=train_matrix, hidden=1,  
          threshold = 0.01,  
          learningrate.limit = NULL,  
          learningrate.factor =  
            list(minus = 0.5, plus = 1.2),  
          algorithm = "rprop+") -> nn1  
plot(nn1)
```



Error: 99.949375 Steps: 23

```
#plug in the test matrix to the NN and name it
output <- compute(nn1, test_matrix[, -c(1,6)], rep=1)
summary(output)
##           Length Class  Mode
## neurons      2    -none- list
## net.result 200    -none- numeric

#create a subset of just the outcomes that we need, which is
#the food desert predictions
output$net.result->pred

#The outcome is many numbers between 0 and 1, so we tell the data to
#change any number above .5 into a 1, and any number below a .5, into a 0
ifelse(output$net.result>.5, 1, 0)->pred

confusionMatrix(as.factor(pred), as.factor(test$foodDesert))

## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0   1
##           0 80 35
##           1 14 71
##
##           Accuracy : 0.755
##           95% CI : (0.6894, 0.8129)
##           No Information Rate : 0.53
```

```
##      P-Value [Acc > NIR] : 4.656e-11
##
##              Kappa : 0.5144
##
## Mcnemar's Test P-Value : 0.004275
##
##      Sensitivity : 0.8511
##      Specificity : 0.6698
##      Pos Pred Value : 0.6957
##      Neg Pred Value : 0.8353
##      Prevalence : 0.4700
##      Detection Rate : 0.4000
##      Detection Prevalence : 0.5750
##      Balanced Accuracy : 0.7604
##
##      'Positive' Class : 0
##
```

*## Must make sure that the data we are testing is the same variable type pred
#is a number and test\$fooddesert is a factor, so we need to include the
#as.factor function for pred in order for confusion matrix to run correctly
#can change the hidden layers, increase the number of neurons, add more data,
#or change the learning algorithm parameters, to try and increase accuracy*

KNN

```
set.seed(341)
```

create a smaller dataset

```
myvars<-c("Urban", "OHU2010", "LowIncomeTracts", "PovertyRate", "MedianFamilyI  
ncome", "foodDesert")
train2<- train[myvars]
test2<- test[myvars]
as.factor(train2$foodDesert)->train2$foodDesert
as.factor(test2$foodDesert)->test2$foodDesert
```

##Used many numbers in tuneLength to help improve accuracy- 5,15,25,27,28,29

```
train(foodDesert~., data=train2, method="knn", tuneLength=28)->yes
```

```
predict(yes, test2)->guess
```

##made both the variables factors

```
confusionMatrix(guess, test2$foodDesert)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##      Reference
```

```
## Prediction 0 1
```

```
##      0 83 38
```

```
##      1 11 68
```

```
##
```

```
##      Accuracy : 0.755
```

```
##      95% CI : (0.6894, 0.8129)
```

```
##      No Information Rate : 0.53
##      P-Value [Acc > NIR] : 4.656e-11
##
##              Kappa : 0.5161
##
##  Mcnemar's Test P-Value : 0.0002038
##
##              Sensitivity : 0.8830
##              Specificity : 0.6415
##              Pos Pred Value : 0.6860
##              Neg Pred Value : 0.8608
##              Prevalence : 0.4700
##              Detection Rate : 0.4150
##      Detection Prevalence : 0.6050
##              Balanced Accuracy : 0.7622
##
##      'Positive' Class : 0
##
```

RANDOM FOREST

```
set.seed(341)
library(randomForest)

## foodDesert must be a factor for random forest to run as a classification
as.factor(train2$foodDesert)->train2$foodDesert
as.factor(test2$foodDesert)->test2$foodDesert
rfm <- randomForest(foodDesert~., data=train2, ntree=300, importance=T)
rfm
##
## Call:
## randomForest(formula = foodDesert ~ ., data = train2, ntree = 300, i
importance = T)
##              Type of random forest: classification
##              Number of trees: 300
## No. of variables tried at each split: 2
##
##      OOB estimate of  error rate: 27.38%
## Confusion matrix:
##      0   1 class.error
## 0 333  73  0.1798030
## 1 146 248  0.3705584

predRF <- predict(rfm, test2, type=c("class"))
confusionMatrix(predRF, test2$foodDesert)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction  0   1
##      0  79  36
```

```
##          1 15 70
##
##          Accuracy : 0.745
##          95% CI : (0.6787, 0.8039)
##    No Information Rate : 0.53
##    P-Value [Acc > NIR] : 3.353e-10
##
##          Kappa : 0.4945
##
##    McNemar's Test P-Value : 0.005101
##
##          Sensitivity : 0.8404
##          Specificity : 0.6604
##          Pos Pred Value : 0.6870
##          Neg Pred Value : 0.8235
##          Prevalence : 0.4700
##          Detection Rate : 0.3950
##          Detection Prevalence : 0.5750
##          Balanced Accuracy : 0.7504
##
##          'Positive' Class : 0
##
```